
Recommender Systems

— Boston University CS506 - Lance Galletti —

An Example: Movie Recommendations

Given

- Users: U_1, \dots, U_n
- Movies: M_1, \dots, M_m
- Ratings: R_{ij}

Goal: Recommend movies to users

Challenges:

- Scale (millions of users, millions of movies)
- Cold Start (change in user base, change in content)
- Sparse Data (Not many users rank movies)

An Example: Movie Recommendations

	\mathbf{M}_1	\mathbf{M}_2	\mathbf{M}_3	\mathbf{M}_4
\mathbf{U}_1	R_{11}	R_{12}	R_{13}	R_{14}
\mathbf{U}_2	R_{21}	R_{22}	R_{23}	R_{24}
\mathbf{U}_3	R_{31}	R_{32}	R_{33}	R_{34}

Use Rating prediction as proxy for recommendation!

An Example: Movie Recommendations

	M_1	M_2	M_3	M_4
U_1	5	5	0	0
U_2	5	4	0	0
U_3	0	0	4	5

An Example: Movie Recommendations

	M_1	M_2	M_3	M_4
U_1	5	?	0	0
U_2	?	4	0	0
U_3	0	?	4	?

Neighborhood Methods

- (user, user) similarity measure
 - i.e. recommend same movies to similar users
- (item, item) similarity measure
 - i.e. recommend movies that are similar

Pros:

- Intuitive / easy to explain
- No training
- Handles new users/items

Challenges:

- Users rate differently (bias)
- Ratings change over time (bias)

Feature Extraction - Content-Based

Suppose we have a set of features that characterizes each movie (ex: category, genre...), we could obtain the following **feature-to-movie** similarity matrix:

	M_1	M_2	M_3	M_4
F_1 (Romance)	.9	1	.1	0
F_2 (Action)	0	.01	1	.9

Feature Extraction - Content-Based

Given this **feature-to-movie** similarity matrix, how can we predict rating for User 2 or Movie 1 (i.e. R_{12})?

If we had a **user-to-feature** similarity matrix, we could multiply:

$$\text{user-to-feature} \times \text{feature-to-movie} = \text{user-to-movie} = R_{ij}$$

Feature Extraction - Content-Based

	F_1 (Romance)	F_2 (Action)
U_1	5	0
U_2	5	0
U_3	0	5

X

	M_1	M_2	M_3	M_4
F_1 (Romance)	.9	1	.1	0
F_2 (Action)	0	.01	1	.9


Feature Extraction - Content-Based

$$\begin{aligned} P^{(2)} &= \begin{bmatrix} 5 \\ 0 \end{bmatrix} & R_{21} &= P^{(2)T} \cdot Q^{(1)} \\ Q^{(1)} &= \begin{bmatrix} .9 \\ 0 \end{bmatrix} & &= \begin{bmatrix} 5 & 0 \end{bmatrix} \cdot \begin{bmatrix} .9 \\ 0 \end{bmatrix} \\ & & &= 4.5 \end{aligned}$$

But, how to we find $p^{(1)}, \dots, p^{(n)}$?

Feature Extraction - Content-Based

$$P^{(j)} = \arg \min_P \frac{1}{\|M^{(j)}\|} \sum_{i \in M^{(j)}} (P^T Q^{(i)} - r_{ij})^2 + \lambda \|p\|^2$$



Regularization Term: a penalty on the size of the parameter p

Feature Extraction

Challenge with content-based:

How to get the right features f_1, \dots, f_k ?

Can we learn these features?

$$\mathbf{R} = \mathbf{PQ}$$

Feature Extraction

Can't use SVD because R is sparse... BUT, we can formulate an optimization problem to solve:

$$\min_{p, q} \sum_{i, j \in R} (r_{ij} - p_i^T q_j)^2 + \lambda (\|p\|_F^2 + \|q\|_F^2)$$

To solve, take derivatives wrt P & Q . Then, just like Expectation-Maximization Algorithm from GMM:

1. Start with random Q
2. Get P
3. Improve Q
4. Repeat 2 & 3